

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
```

# Machine Learning Algorithms:

- Decision Tree
- Logistic Regression

Experiment 1: Data Splitting with the hybrid oversampling and undersampling- In the first experiment, we split the dataset into training (80%) and testing (20%) subsets, maintaining class balance through stratification.

Experiment 2: Stratified Cross-Validation with SMOTE (Oversampling)  
For the third experiment, we enhance our stratified cross-validation by incorporating Synthetic Minority Over-sampling Technique (SMOTE). SMOTE is used to address class\ imbalance by generating synthetic samples, improving the model's ability to learn from the minority class.

```
In [2]: data = pd.read_csv('credit_card_churn.csv')
```

```
In [3]: data.head(2)
```


Out[3]:

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	Education_Level	Marital_Status	Income_Category	Card_Category	Months_on_book	...	Months_Inactive_12_mon	Contacts_Count_12_mon	Credit_Limit	Total_Revo
0	768805383	Existing Customer	45	M	3	High School	Married	60K–80K	Blue	39	...	1	3	12691.0	
1	818770008	Existing Customer	49	F	5	Graduate	Single	Less than \$40K	Blue	44	...	1	2	8256.0	

2 rows × 21 columns

```
In [4]: data['Attrition_Flag'].value_counts().plot.pie(autopct='%0.2f')
#The target variable is not balance
```

```
Out[4]: <AxesSubplot:ylabel='Attrition_Flag'>
```



```
In [5]: Finaldata = data[['Attrition_Flag', 'Gender', 'Income_Category', 'Total_Trans_Ct', 'Avg_Utilization_Ratio', 'Total_Revolving_Bal', 'Months_Inactive_12_mon', 'Total_Trans_Amt', 'Total_Amt_Chng
```

In [6]: Finaldata.dtypes

Out[6]: Attrition\_Flag object  
Gender object  
Income\_Category object  
Total\_Trans\_Ct int64  
Avg\_Utilization\_Ratio float64  
Total\_Revolving\_Bal int64  
Months\_Inactive\_12\_mon int64  
Total\_Trans\_Amt int64  
Total\_Amt\_Chng\_Q4\_Q1 float64  
dtype: object

In [7]: from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()  
Finaldata['Gender'] = le.fit\_transform(data['Gender'])  
Finaldata['Income\_Category'] = le.fit\_transform(data['Income\_Category'])  
Finaldata['Attrition\_Flag'] = le.fit\_transform(data['Attrition\_Flag'])

C:\Users\Romelio Villar Jr\AppData\Local\Temp\ipykernel\_6900\320870704.py:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))  
Finaldata['Gender'] = le.fit\_transform(data['Gender'])  
C:\Users\Romelio Villar Jr\AppData\Local\Temp\ipykernel\_6900\320870704.py:4: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))  
Finaldata['Income\_Category'] = le.fit\_transform(data['Income\_Category'])  
C:\Users\Romelio Villar Jr\AppData\Local\Temp\ipykernel\_6900\320870704.py:5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))  
Finaldata['Attrition\_Flag'] = le.fit\_transform(data['Attrition\_Flag'])

In [8]: Finaldata.head(2)

Out[8]:

	Attrition_Flag	Gender	Income_Category	Total_Trans_Ct	Avg_Utilization_Ratio	Total_Revolving_Bal	Months_Inactive_12_mon	Total_Trans_Amt	Total_Amt_Chng_Q4_Q1
0	1	1	2	42	0.061	777	1	1144	1.335
1	1	0	4	33	0.105	864	1	1291	1.541

In [9]: Finaldata.dtypes

Out[9]: Attrition\_Flag int32  
Gender int32  
Income\_Category int32  
Total\_Trans\_Ct int64  
Avg\_Utilization\_Ratio float64  
Total\_Revolving\_Bal int64  
Months\_Inactive\_12\_mon int64  
Total\_Trans\_Amt int64  
Total\_Amt\_Chng\_Q4\_Q1 float64  
dtype: object

In [10]: X = Finaldata.drop('Attrition\_Flag', axis = 1)  
y = Finaldata['Attrition\_Flag']

Experiment 1: Data Splitting training (80%) and testing (20%) subsets, stratify=yes, hyperparameters

```
In [11]: import numpy as np
from sklearn.model_selection import StratifiedKFold, GridSearchCV, train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, roc_auc_score, precision_score, recall_score, f1_score, confusion_matrix
from imblearn.over_sampling import RandomOverSampler, SMOTE
from imblearn.under_sampling import RandomUnderSampler
```

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
In [13]: print(y_train.value_counts(normalize=True)*100)
```

```
1    83.92791
0    16.07209
Name: Attrition_Flag, dtype: float64
```

```
In [14]: print(y_test.value_counts(normalize=True)*100)
```

```
1    83.958539
0    16.041461
Name: Attrition_Flag, dtype: float64
```

## Decision Tree

Experiment 1

```
In [15]: dt_classifier = DecisionTreeClassifier(random_state=42)
```

```
In [16]: param_grid_dt = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', 'log2']
}
```

```
In [17]: clf = GridSearchCV(dt_classifier, param_grid_dt)
clf
```

```
Out[17]: GridSearchCV(estimator=DecisionTreeClassifier(random_state=42),
                      param_grid={'criterion': ['gini', 'entropy'],
                                   'max_depth': [None, 10, 20, 30],
                                   'max_features': ['auto', 'sqrt', 'log2'],
                                   'min_samples_leaf': [1, 2, 4],
                                   'min_samples_split': [2, 5, 10]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\tree\_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the
past behaviour, explicitly set `max_features='sqrt'`.
  warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\tree\_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the
past behaviour, explicitly set `max_features='sqrt'`.
  warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\tree\_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the
past behaviour, explicitly set `max_features='sqrt'`.
  warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\tree\_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the
past behaviour, explicitly set `max_features='sqrt'`.
  warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\tree\_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the
past behaviour, explicitly set `max_features='sqrt'`.
  warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\tree\_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the
past behaviour, explicitly set `max_features='sqrt'`.
  warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\tree\_classes.py:269: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the
past behaviour, explicitly set `max_features='sqrt'`.
```

**Out[19]:** DecisionTreeClassifier(max\_depth=10, max\_features='log2', min\_samples\_leaf=2, min\_samples\_split=10, random\_state=42)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook. On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

**Out[21]:** DecisionTreeClassifier(max\_depth=10, max\_features='log2', min\_samples\_leaf=2, min\_samples\_split=10, random\_state=42)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [26]: accuracy = accuracy_score(y_test, y_pred_d1)
precision = precision_score(y_test, y_pred_d1)
recall = recall_score(y_test,y_pred_d1)
f1 = f1_score(y_test, y_pred_d1)
confusion = confusion_matrix(y_test, y_pred_d1)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("f1_score:", f1_score)
print("confusion matrix\n", confusion)
```

```
Accuracy: 0.914116485686081
Precision: 0.9607724803862402
Recall: 0.9359200470311582
f1_score: <function f1_score at 0x000001C107307CA0>
confusion matrix
[[ 260   65]
 [ 109 1592]]
```

Undersampling

```
In [27]: # Now, apply RandomOverSampler to the training data
rus = RandomUnderSampler(sampling_strategy=0.5) # Adjust the sampling strategy as needed
X_train_resampled_U, y_train_resampled_U = rus.fit_resample(X_train, y_train)
```

```
In [28]: best_dt_classifier.fit(X_train_resampled_U, y_train_resampled_U)
```

```
Out[28]: DecisionTreeClassifier(max_depth=10, max_features='log2', min_samples_leaf=2,
                                min_samples_split=10, random_state=42)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [29]: y_pred_d2 = best_dt_classifier.predict(X_test)
```

```
In [31]: accuracy = accuracy_score(y_test, y_pred_d2)
precision = precision_score(y_test, y_pred_d2)
recall = recall_score(y_test,y_pred_d2)
f1 = f1_score(y_test, y_pred_d2)
confusion = confusion_matrix(y_test, y_pred_d2)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("f1_score:", f1_score)
print("confusion matrix\n", confusion)
```

```
Accuracy: 0.9106614017769002
Precision: 0.9679802955665024
Recall: 0.9241622574955908
f1_score: <function f1_score at 0x000001C107307CA0>
confusion matrix
[[ 273   52]
 [ 129 1572]]
```

Experiment 2: Decision Tree with Stratified Cross Validation

```
In [32]: cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```



```
In [45]: cv_metrics_R = cross_val_score(best_dt_classifier, X_train_smote, y_train_smote, cv=cv, scoring='recall')
cv_metrics_R
```

```
Out[45]: array([0.92573529, 0.91617647, 0.90955882, 0.93230316, 0.89485294])
```

```
In [46]: print('mean_precision:', np.mean(cv_metrics_P))
print('mean_recall:', np.mean(cv_metrics_R))
```

```
mean_precision: 0.9316167639960244
mean_recall: 0.9157253387006016
```

## Logistic Regression

Experiment 1

```
In [47]: LR_classifier = LogisticRegression()
```

```
In [48]: param_grid_LR = {
    'penalty' : ['l1', 'l2'],
    'C' : [0.001, 0.01, 0.1, 1.0],
    'solver' : ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'],
    'max_iter':[100, 200, 300]
}
```

```
In [49]: R_clf = GridSearchCV(LR_classifier, param_grid_LR)
```

```
In [50]: R_clf.fit(X_train, y_train)
```

```
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
In [51]: LR_best_parameter = R_clf.best_estimator_
LR_best_parameter
```

```
Out[51]: LogisticRegression(solver='newton-cg')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [52]: LR_ros = RandomOverSampler(sampling_strategy = 'auto')
X_train_L, y_train_L = LR_ros.fit_resample(X_train, y_train)
```



```
In [53]: LR_best_parameter.fit(X_train_L, y_train_L)
```

```
Out[53]: LogisticRegression(solver='newton-cg')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [54]: y_pred_LR = LR_best_parameter.predict(X_test)
```

```
In [56]: accuracy = accuracy_score(y_test, y_pred_LR)
precision = precision_score(y_test, y_pred_LR)
recall = recall_score(y_test,y_pred_LR)
f1 = f1_score(y_test, y_pred_LR)
confusion = confusion_matrix(y_test, y_pred_LR)
```

```
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("f1_score:", f1_score)
print("confusion matrix\n", confusion)
```

```
Accuracy: 0.7966436327739388
Precision: 0.9593727726300784
Recall: 0.7912992357436802
f1_score: <function f1_score at 0x000001C107307CA0>
confusion matrix
[[ 268   57]
 [ 355 1346]]
```

Undersampling

```
In [57]: LR_rus = RandomUnderSampler(sampling_strategy = 'auto')
X_train_LR, y_train_LR = LR_rus.fit_resample(X_train, y_train)
```

```
In [58]: LR_best_parameter.fit(X_train_LR, y_train_LR)
```

```
Out[58]: LogisticRegression(solver='newton-cg')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [59]: y_pred_LR2 = LR_best_parameter.predict(X_test)
```

```
In [61]: accuracy = accuracy_score(y_test, y_pred_LR2)
precision = precision_score(y_test, y_pred_LR2)
recall = recall_score(y_test,y_pred_LR2)
f1 = f1_score(y_test, y_pred_LR2)
confusion = confusion_matrix(y_test, y_pred_LR2)
```

```
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("f1_score:", f1_score)
print("confusion matrix\n", confusion)
```

```
Accuracy: 0.7951628825271471
Precision: 0.9599427753934192
Recall: 0.7889476778365667
f1_score: <function f1_score at 0x000001C107307CA0>
confusion matrix
[[ 269   56]
 [ 359 1342]]
```



## Experiment 2

```
In [62]: smote_LR = SMOTE(sampling_strategy='auto')
```

```
In [63]: R_clf = GridSearchCV(LR_classifier, param_grid_LR, cv=cv)
```

```
In [64]: R_clf.fit(X_train, y_train)
```

```
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\scipy\optimize\_linesearch.py:416: LineSearchWarning: Rounding errors prevent the line search from converging
warn(msg, LineSearchWarning)
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\scipy\optimize\_linesearch.py:306: LineSearchWarning: The line search algorithm did not converge
warn('The line search algorithm did not converge', LineSearchWarning)
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\utils\optimize.py:203: UserWarning: Line Search failed
warnings.warn("Line Search failed")
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\utils\optimize.py:210: ConvergenceWarning: newton-cg failed to converge. Increase the number of iterations.
warnings.warn(
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT
```

```
In [117]: X_train_LRsmote, y_train_LRsmote = smote_LR.fit_resample(X_train, y_train)
```

```
In [118]: LR_best_parameter.fit(X_train_LRsmote, y_train_LRsmote)
```

```
Out[118]: LogisticRegression(solver='newton-cg')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**  
**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [119]: y_pred_smlr = LR_best_parameter.predict(X_test)
```

```
In [114]: LR_metrics_P = cross_val_score(LR_best_parameter, X_train_LRsmote, y_train_LRsmote, cv=cv, scoring='precision')
LR_metrics_P
```

```
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\scipy\optimize\_linesearch.py:306: LineSearchWarning: The line search algorithm did not converge
warn('The line search algorithm did not converge', LineSearchWarning)
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\utils\optimize.py:203: UserWarning: Line Search failed
warnings.warn("Line Search failed")
```

```
Out[114]: array([0.8515219 , 0.86712225, 0.87372549, 0.8453997 , 0.85352761])
```

```
In [115]: LR_metrics_R = cross_val_score(LR_best_parameter, X_train_LRsmote, y_train_LRsmote, cv=cv, scoring='recall')
LR_metrics_R
```

```
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\scipy\optimize\_linesearch.py:306: LineSearchWarning: The line search algorithm did not converge
warn('The line search algorithm did not converge', LineSearchWarning)
C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\sklearn\utils\optimize.py:203: UserWarning: Line Search failed
warnings.warn("Line Search failed")
```

```
Out[115]: array([0.84338235, 0.83970588, 0.81911765, 0.82487123, 0.81838235])
```

```
In [116]: print('Precision_Logistic Regression:', np.mean(LR_metrics_P))  
          print('Recall_Logistic Regression:', np.mean(LR_metrics_R))
```

```
Precision_Logistic Regression: 0.8582593887901722  
Recall_Logistic Regression: 0.8290918928277712
```

```
In [ ]:
```