

```
In [60]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats

In [61]: data = pd.read_csv('credit_card_churn.csv')
data.head(2)

Out[61]:
```

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	Education_Level	Marital_Status	Income_Category	Card_Category	Months_on_book	...
0	768805383	Existing Customer	45	M	3	High School	Married	60K–80K	Blue	39	...
1	818770008	Existing Customer	49	F	5	Graduate	Single	Less than \$40K	Blue	44	...

2 rows × 21 columns

Checking nulls and duplicates

```
In [3]: data.isnull().sum()

Out[3]: CLIENTNUM      0
Attrition_Flag      0
Customer_Age        0
Gender              0
Dependent_count     0
Education_Level     0
Marital_Status      0
Income_Category     0
Card_Category       0
Months_on_book      0
Total_Relationship_Count  0
Months_Inactive_12_mon  0
Contacts_Count_12_mon  0
Credit_Limit        0
Total_Revolving_Bal  0
Avg_Open_To_Buy     0
Total_Amt_Chng_Q4_Q1  0
Total_Trans_Amt      0
Total_Trans_Ct       0
Total_Ct_Chng_Q4_Q1  0
Avg_Utilization_Ratio  0
dtype: int64

In [4]: data.duplicated().sum()

Out[4]: 0

In [5]: data.dtypes

Out[5]: CLIENTNUM      int64
Attrition_Flag    object
Customer_Age      int64
Gender            object
Dependent_count   int64
Education_Level   object
Marital_Status    object
Income_Category   object
Card_Category     object
Months_on_book    int64
Total_Relationship_Count  int64
Months_Inactive_12_mon  int64
Contacts_Count_12_mon  int64
Credit_Limit      float64
Total_Revolving_Bal  int64
Avg_Open_To_Buy    float64
Total_Amt_Chng_Q4_Q1  float64
Total_Trans_Amt    int64
Total_Trans_Ct     int64
Total_Ct_Chng_Q4_Q1  float64
Avg_Utilization_Ratio  float64
dtype: object
```

Summary Statistics

In [6]:

data.describe()

Out[6]:

	CLIENTNUM	Customer_Age	Dependent_count	Months_on_book	Total_Relationship_Count	Months_Inactive_12_mon	Contacts_Count_12_mon	Credit_Lim
count	1.012700e+04	10127.000000	10127.000000	10127.000000	10127.000000	10127.000000	10127.000000	10127.000000
mean	7.391776e+08	46.325960	2.346203	35.928409	3.812580	2.341167	2.455317	8631.95369
std	3.690378e+07	8.016814	1.298908	7.986416	1.554408	1.010622	1.106225	9088.77669
min	7.080821e+08	26.000000	0.000000	13.000000	1.000000	0.000000	0.000000	1438.30000
25%	7.130368e+08	41.000000	1.000000	31.000000	3.000000	2.000000	2.000000	2555.00000
50%	7.179264e+08	46.000000	2.000000	36.000000	4.000000	2.000000	2.000000	4549.00000
75%	7.731435e+08	52.000000	3.000000	40.000000	5.000000	3.000000	3.000000	11067.50000
max	8.283431e+08	73.000000	5.000000	56.000000	6.000000	6.000000	6.000000	34516.00000

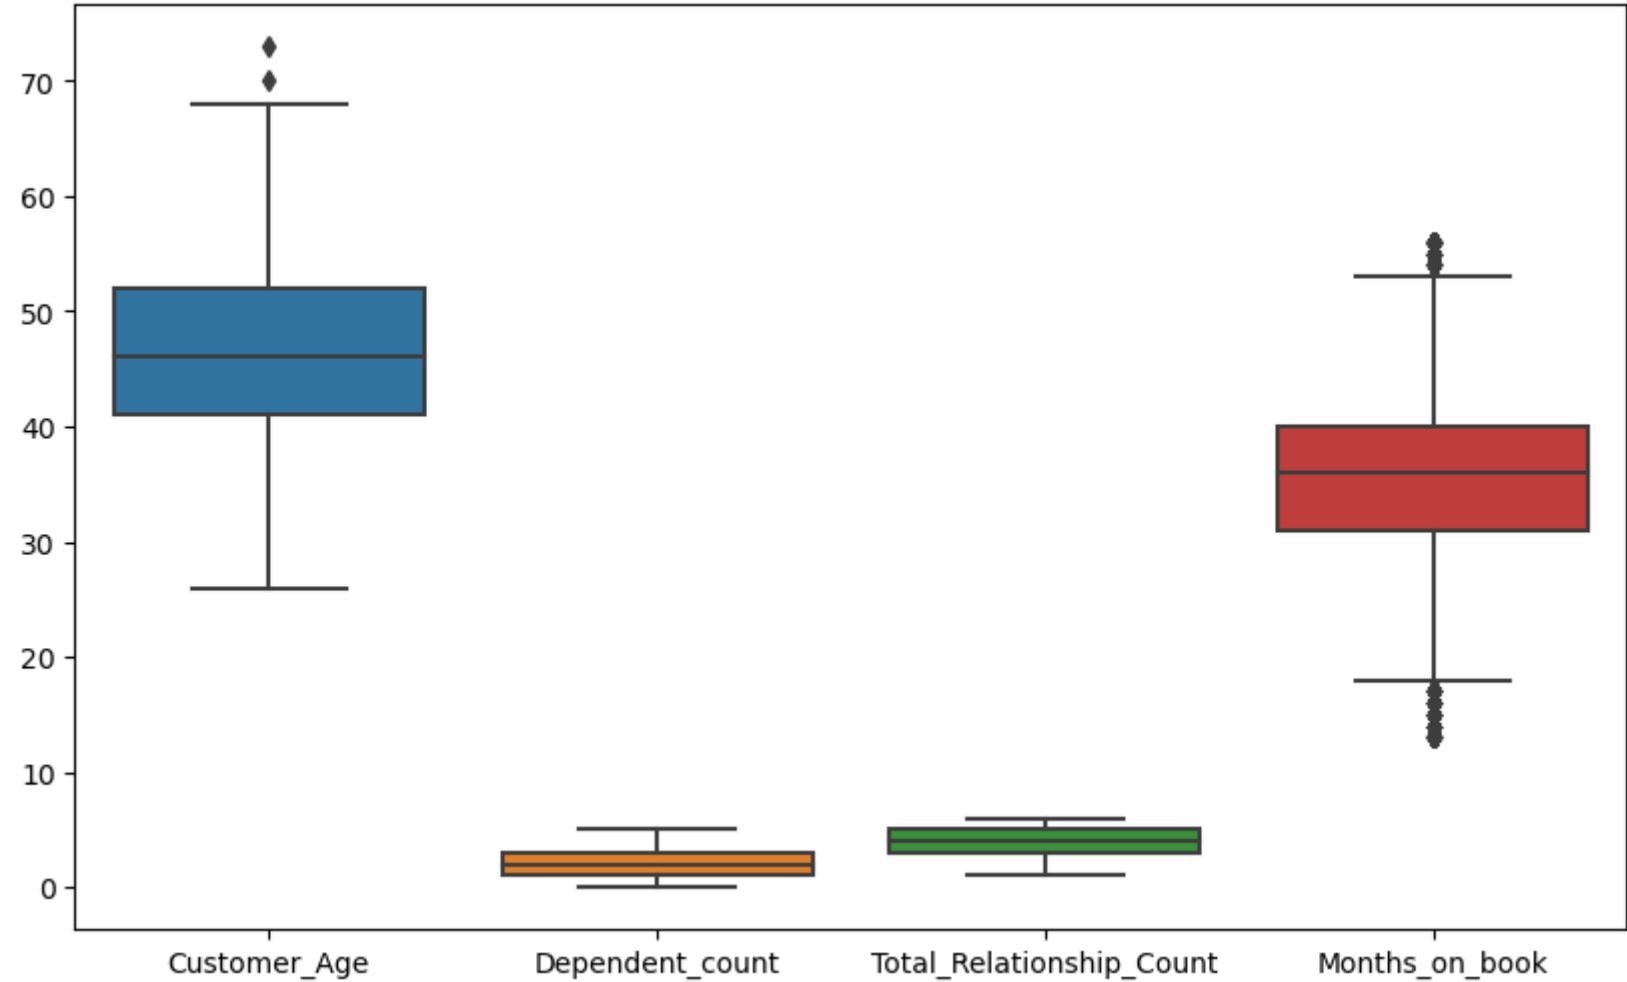
Checking Outliers

In [7]:

data_number1= data[['Customer_Age', 'Dependent_count', 'Total_Relationship_Count','Months_on_book']]

In [8]:

plt.figure(figsize=(10,6))
sns.boxplot(data= data_number1)
plt.show()

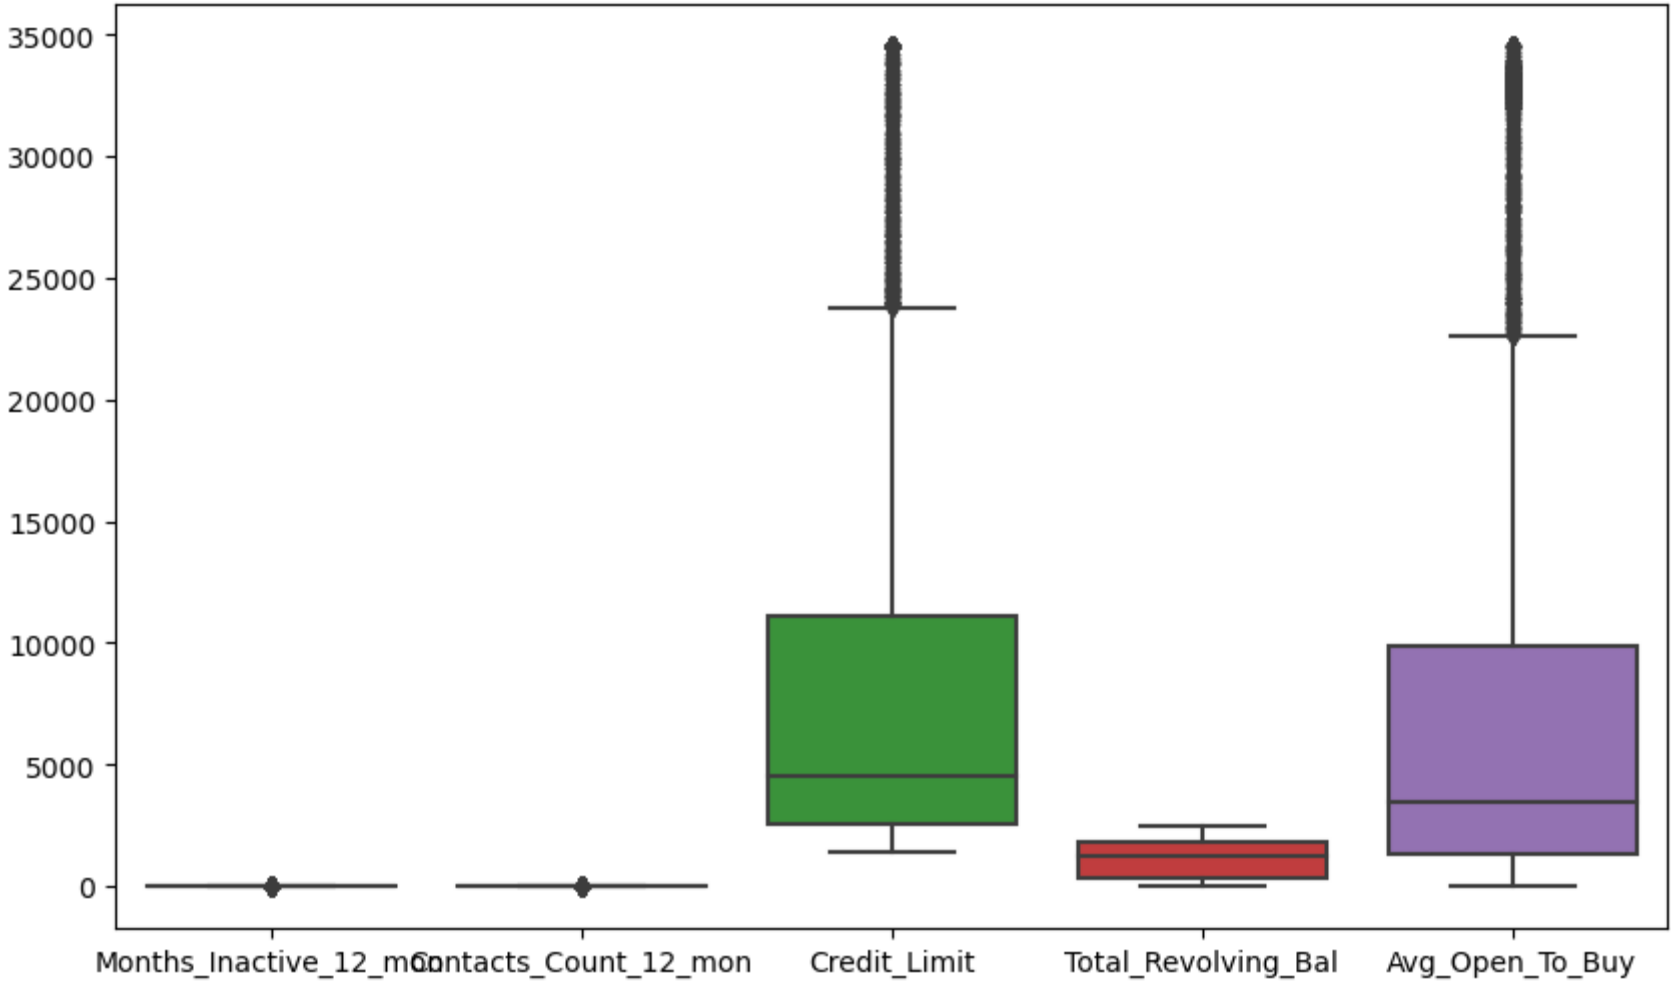


In [9]:

data_number2 = data[['Months_Inactive_12_mon','Contacts_Count_12_mon', 'Credit_Limit', 'Total_Revolving_Bal', 'Avg_Open_To_Buy']]

```
In [10]: plt.figure(figsize=(10,6))
sns.boxplot(data=data_number2)
```


Out[10]: <AxesSubplot:>



```
In [11]: data_number3 = data[['Total_Amt_Chng_Q4_Q1', 'Total_Trans_Amt', 'Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1', 'Avg_Utilization_Ratio']]
```


```
In [12]: plt.figure(figsize=(10,12))
sns.boxplot(data=data_number3)
```

Out[12]: <AxesSubplot:>



```
In [13]: data['Attrition_Flag'].value_counts().plot.pie(autopct='%.2f')
#The data is not balance
```

Out[13]: <AxesSubplot:ylabel='Attrition_Flag'>




```
In [112]: counts = data.groupby(['Attrition_Flag', 'Gender']).size().unstack(fill_value=0)
counts
```

Out[112]:

	Gender	F	M
Attrition_Flag			
Attrited Customer		930	697
Existing Customer		4428	4072

```
In [113]: counts.plot(kind='barh', stacked=True)
```

Out[113]: <AxesSubplot:ylabel='Attrition_Flag'>




```
In [114]: counts_2 = data.groupby(['Attrition_Flag', 'Income_Category']).size().unstack(fill_value=0)
counts_2
```

Out[114]:

Income_Category	\$120K +	40K – 60K	60K – 80K	80K – 120K	120K – 240K	Less than \$40K	Unknown
Attrition_Flag							
Attrited Customer	126	271	189	242	612	187	
Existing Customer	601	1519	1213	1293	2949	925	

```
In [116]: counts_2.plot(kind='barh', stacked=True)
```

Out[116]: <AxesSubplot:ylabel='Attrition_Flag'>



```
In [14]: from sklearn.preprocessing import LabelEncoder
```

```
In [15]: le = LabelEncoder()
data['Attrition_Flag'] = le.fit_transform(data['Attrition_Flag'])
data.head(2)

Out[15]:
```

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	Education_Level	Marital_Status	Income_Category	Card_Category	Months_on_book	...
0	768805383	1	45	M	3	High School	Married	60K–80K	Blue	39	...
1	818770008	1	49	F	5	Graduate	Single	Less than \$40K	Blue	44	...

2 rows × 21 columns

Checking the correlation

```
In [16]: data_corr = data.corr()

In [17]: plt.figure(figsize=(14,12))
sns.heatmap(data_corr, cbar=True, annot=True)

Out[17]: <AxesSubplot:>
```



Checking the skewness

```
In [18]: from scipy.stats import skew

In [19]: Num_data = data[['Customer_Age', 'Dependent_count', 'Months_on_book',
'Total_Relationship_Count',
'Months_Inactive_12_mon',
'Contacts_Count_12_mon',
'Credit_Limit',
'Total_Revolving_Bal',
'Avg_Open_To_Buy',
'Total_Amt_Chng_Q4_Q1',
'Total_Trans_Amt',
'Total_Trans_Ct',
'Total_Ct_Chng_Q4_Q1',
'Avg_Utilization_Ratio']]

In [20]: for col in Num_data:
print(col)
print(skew(Num_data[col]))

plt.figure()
sns.distplot(Num_data[col])

Customer_Age
-0.03360003857464426

Dependent_count
-0.02082245083419453
Months_on_book
-0.1065495749017217

C:\Users\Romelio Villar Jr\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated f
unction and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with sim
ilar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```

Feature Selection

```
In [ ]:
```

We will use Point Biserial. The point biserial correlation coefficient, which is a special case of Pearson's correlation coefficient. It measures the relationship between two variables
One continuous variable and One naturally binary variable.
In our dataset, the "ATTRITION FLAG" is a binary variable.
We will examine the correlation of numerical variables with the "Attrition_flag".

```
In [21]: data['Attrition_Flag'].dtypes
```

```
Out[21]: dtype('int32')
```

```
In [22]: target = data['Attrition_Flag']
feature = data[['Customer_Age', 'Dependent_count', 'Months_on_book',
'Total_Relationship_Count',
'Months_Inactive_12_mon',
'Contacts_Count_12_mon',
'Credit_Limit',
'Total_Revolving_Bal',
'Avg_Open_To_Buy',
'Total_Amt_Chng_Q4_Q1',
'Total_Trans_Amt',
'Total_Trans_Ct',
'Total_Ct_Chng_Q4_Q1',
'Avg_Utilization_Ratio']]
```

```
In [23]: from sklearn.preprocessing import LabelEncoder
```

```
In [24]: le = LabelEncoder()
```

```
In [25]: data['Attrition_Flag'] = le.fit_transform(target)
data['Attrition_Flag']
```

```
Out[25]: 0         1
1         1
2         1
3         1
4         1
..
10122     1
10123     0
10124     0
10125     0
10126     0
Name: Attrition_Flag, Length: 10127, dtype: int64
```

```
In [26]: feature.astype('int32').dtypes
```

```
Out[26]: Customer_Age           int32
Dependent_count           int32
Months_on_book            int32
Total_Relationship_Count  int32
Months_Inactive_12_mon    int32
Contacts_Count_12_mon     int32
Credit_Limit              int32
Total_Revolving_Bal       int32
Avg_Open_To_Buy           int32
Total_Amt_Chng_Q4_Q1      int32
Total_Trans_Amt           int32
Total_Trans_Ct            int32
Total_Ct_Chng_Q4_Q1       int32
Avg_Utilization_Ratio     int32
dtype: object
```

```
In [27]: print("Target Shape:", target.shape)
print("Feature Shape:", feature.shape)
```

```
Target Shape: (10127,)
Feature Shape: (10127, 14)
```

In [28]:

```
for column in feature.columns:
    point_biserial_corr, p_value = stats.pointbiserialr(target, feature[column])
    print(f'Feature: {column}')
    print(f'Point-Biserial Correlation: {point_biserial_corr}')
    print(f'P-value: {p_value}')
```

Feature: Customer_Age
Point-Biserial Correlation: -0.01820313853255065
P-value: 0.06698688501759016
Feature: Dependent_count
Point-Biserial Correlation: -0.018990596311193708
P-value: 0.056002392535092434
Feature: Months_on_book
Point-Biserial Correlation: -0.01368685117790971
P-value: 0.1684370287649442
Feature: Total_Relationship_Count
Point-Biserial Correlation: 0.15000522801913754
P-value: 4.829281002183993e-52
Feature: Months_Inactive_12_mon
Point-Biserial Correlation: -0.152448806326925
P-value: 1.0326639995930894e-53
Feature: Contacts_Count_12_mon
Point-Biserial Correlation: -0.2044905099816044
P-value: 4.697489630751521e-96
Feature: Credit_Limit
Point-Biserial Correlation: 0.023872994836161524
P-value: 0.01628535720539447
Feature: Total_Revolving_Bal
Point-Biserial Correlation: 0.2630528831292032
P-value: 6.630148455417239e-160
Feature: Avg_Open_To_Buy
Point-Biserial Correlation: 0.00028507749393779595
P-value: 0.977116089445888
Feature: Total_Amt_Chng_Q4_Q1
Point-Biserial Correlation: 0.13106284781447014
P-value: 4.836642703584966e-40
Feature: Total_Trans_Amt
Point-Biserial Correlation: 0.168598381410079
P-value: 1.857438655661277e-65
Feature: Total_Trans_Ct
Point-Biserial Correlation: 0.37140270118892776
P-value: 0.0
Feature: Total_Ct_Chng_Q4_Q1
Point-Biserial Correlation: 0.29005400688089117
P-value: 1.6477247846937629e-195
Feature: Avg_Utilization_Ratio
Point-Biserial Correlation: 0.178410331561747
P-value: 3.3576893282456845e-73

In [29]:

```
data
```

Out[29]:

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	Education_Level	Marital_Status	Income_Category	Card_Category	Months_on_book
0	768805383	1	45	M	3	High School	Married	60K–80K	Blue	39
1	818770008	1	49	F	5	Graduate	Single	Less than \$40K	Blue	44
2	713982108	1	51	M	3	Graduate	Married	80K–120K	Blue	36
3	769911858	1	40	F	4	High School	Unknown	Less than \$40K	Blue	34
4	709106358	1	40	M	3	Uneducated	Married	60K–80K	Blue	21
...
10122	772366833	1	50	M	2	Graduate	Single	40K–60K	Blue	40
10123	710638233	0	41	M	2	Unknown	Divorced	40K–60K	Blue	25
10124	716506083	0	44	F	1	High School	Married	Less than \$40K	Blue	36
10125	717406983	0	30	M	2	Graduate	Unknown	40K–60K	Blue	36
10126	714337233	0	43	F	2	Graduate	Married	Less than \$40K	Silver	25

10127 rows × 11 columns


```
In [30]: from sklearn.preprocessing import LabelEncoder as le

le = LabelEncoder()
data['Marital_Status'] = le.fit_transform(data['Marital_Status'])

le = LabelEncoder()
data['Education_Level'] = le.fit_transform(data['Education_Level'])

le = LabelEncoder()
data['Card_Category'] = le.fit_transform(data['Card_Category'])

# Encode 'Gender'
data['Gender'] = le.fit_transform(data['Gender'])

# Encode 'Income_Category'
data['Income_Category'] = le.fit_transform(data['Income_Category'])
data.head(2)
```

Out[30]:

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	Education_Level	Marital_Status	Income_Category	Card_Category	Months_on_book	...
0	768805383	1	45	1	3	3	1	2	0	39	...
1	818770008	1	49	0	5	2	2	4	0	44	...

2 rows × 21 columns

```
In [31]: data.columns
```

Out[31]:

```
Index(['CLIENTNUM', 'Attrition_Flag', 'Customer_Age', 'Gender',
      'Dependent_count', 'Education_Level', 'Marital_Status',
      'Income_Category', 'Card_Category', 'Months_on_book',
      'Total_Relationship_Count', 'Months_Inactive_12_mon',
      'Contacts_Count_12_mon', 'Credit_Limit', 'Total_Revolving_Bal',
      'Avg_Open_To_Buy', 'Total_Amt_Chng_Q4_Q1', 'Total_Trans_Amt',
      'Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1', 'Avg_Utilization_Ratio'],
      dtype='object')
```

```
In [32]: target.dtypes
```

Out[32]:

```
dtype('int32')
```

```
In [33]: categorical_fetaures = data[['Gender', 'Education_Level', 'Marital_Status', 'Income_Category', 'Card_Category']]
```

```
In [34]: from sklearn.feature_selection import chi2
```

```
In [35]: Chi_scores = chi2(categorical_fetaures, target)
```

```
In [36]: Chi_scores[0]
```

Out[36]:

```
array([7.4432227 , 0.33923111, 1.30275451, 2.47516959, 0.98612022])
```

```
In [37]: Chi_scores[1]
```


Out[37]:

```
array([0.00636758, 0.56027338, 0.25371069, 0.11565697, 0.32069248])
```

```
In [38]: Chi_values = pd.Series(Chi_scores[0], categorical_fetaures.columns)
Chi_values.sort_values(ascending=False, inplace=True)
Chi_values.plot(kind='barh') #the higher the better
```

Out[38]:

```
<AxesSubplot:>
```




```
In [39]: P_values = pd.Series(Chi_scores[1], categorical_fetaures.columns)
P_values.sort_values(ascending=False, inplace=True)
P_values
```

Out[39]: Education_Level 0.560273
Card_Category 0.320692
Marital_Status 0.253711
Income_Category 0.115657
Gender 0.006368
dtype: float64

```
In [40]: P_values.plot(kind='barh', color='red') #The Lower the better
```

Out[40]: <AxesSubplot:>



```
In [41]: Finaldata = data[['Attrition_Flag', 'Gender', 'Income_Category', 'Total_Trans_Ct', 'Avg_Utilization_Ratio', 'Total_Revolving_Bal',
```

```
In [42]: X = Finaldata.drop('Attrition_Flag', axis=1)
y = Finaldata['Attrition_Flag']
```

```
In [49]: data.head(2)
```

Out[49]:

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	Education_Level	Marital_Status	Income_Category	Card_Category	Months_on_book	...
0	768805383	1	45	1	3	3	1	2	0	39	...
1	818770008	1	49	0	5	2	2	4	0	44	...

2 rows × 21 columns